

# Deep Learning Basics

## Recurrent Neural Networks

Francis Steen and Xinyu You  
Red Hen Lab  
August 2019

# A Question

You should take your \_\_\_\_?

What is the most likely word in the blank?

# A Question

It's raining outside. You should take your \_\_\_\_?\_\_\_\_\_.

What is the most likely word in the blank?

# A Question

It's raining outside. You should take your umbrella.

How?

# A Question

It's raining outside. You should take your umbrella.



information

We need information from the distant past  
to accurately predict the correct word.



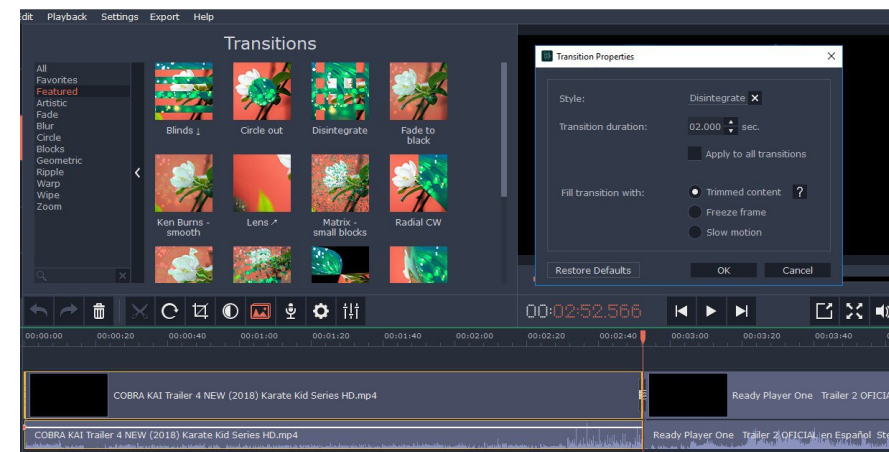
# Sequence Data

- Audio



- Text

- Video



- DAN Sequence

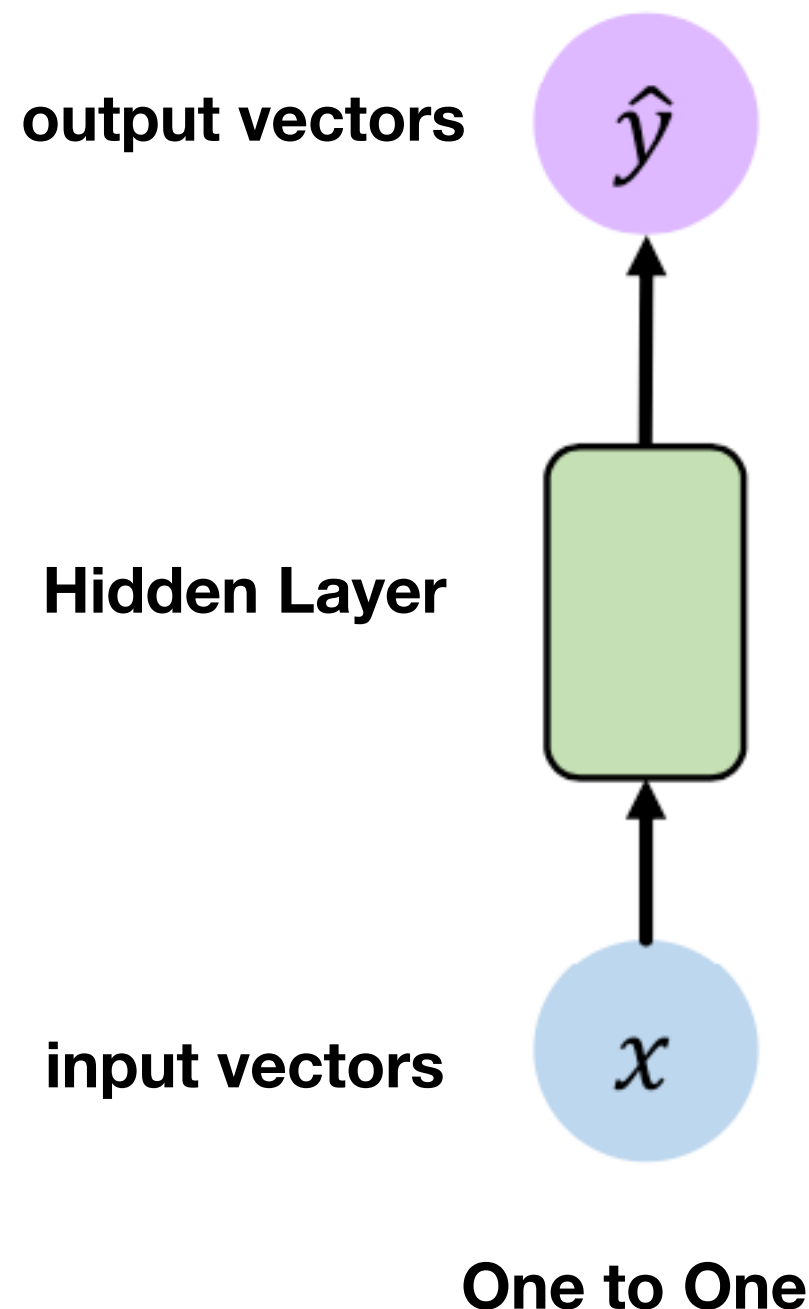
Target Genome	ATTGCGCAGAGACCTAAGGCATTAGCTTGGCCCTAAAG			
Reads	ATTGCG	AGAGACCTAAG	TTAGCTTGGC	AAG
	TGCGCAGA		TGGCCCTAA	
Overlapping	ATTGCG	AGAGACCTAAG	TTAGCTTGGC	AAG
	TGCGCAGA		TGGCCCTAA	
Contigs	ATTGCGCAGAGACCTAAG		TTAGCTTGGCCCTAAAG	

# Sequence modeling

To model sequences, we need to:

1. Handle variable-length sequences
2. Track long-term dependencies
3. Maintain information about order
4. Share parameters across the sequence

# Deep forward neural networks



## One to One

- $x$  is the input vectors which are in blue.
- $y$  is the output vectors which are in purple.
- And green vectors hold the hidden layer.
- Assuming that each input is **independent**, that is, the output of each network depends only on the current input.



# Recurrent neural networks

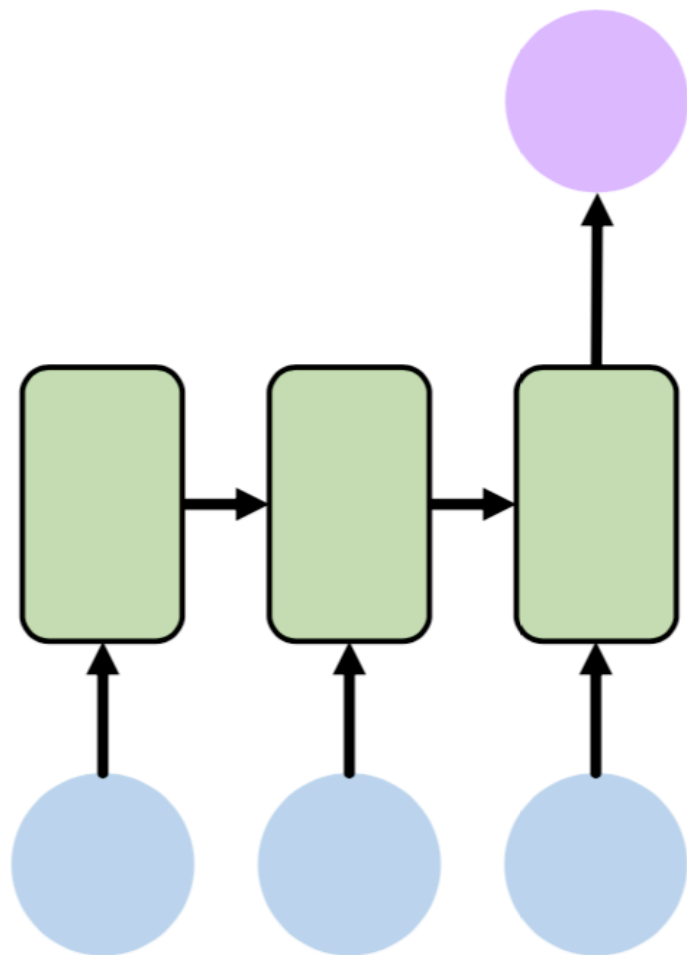
## Many to One

For example:

Input a sequence of words and predict the sentiment associated with it.

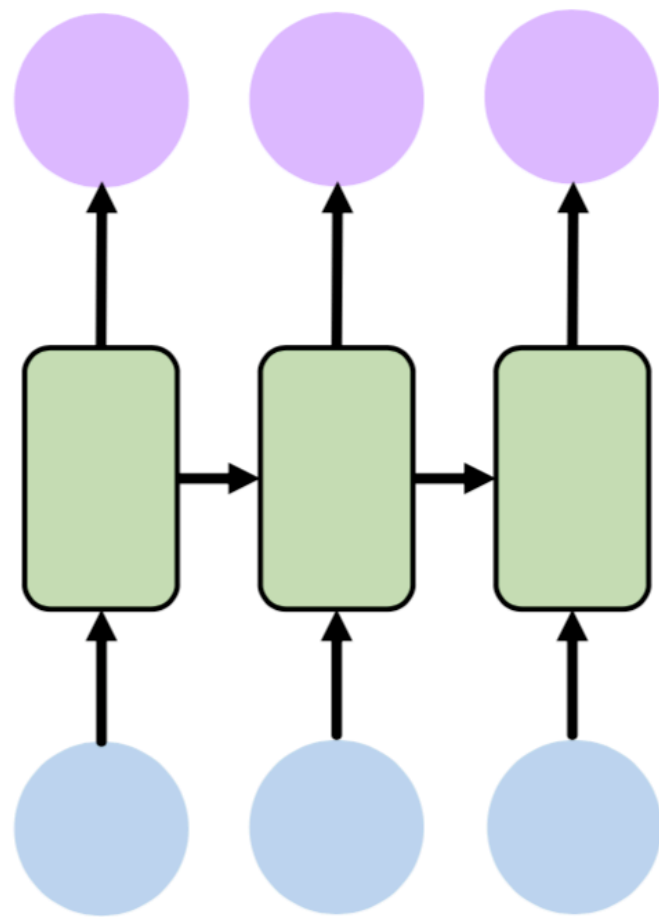
Input: the summer holiday begins, I am so happy.

Output: Happy.



Many to One

# Recurrent neural networks



**Many to Many**

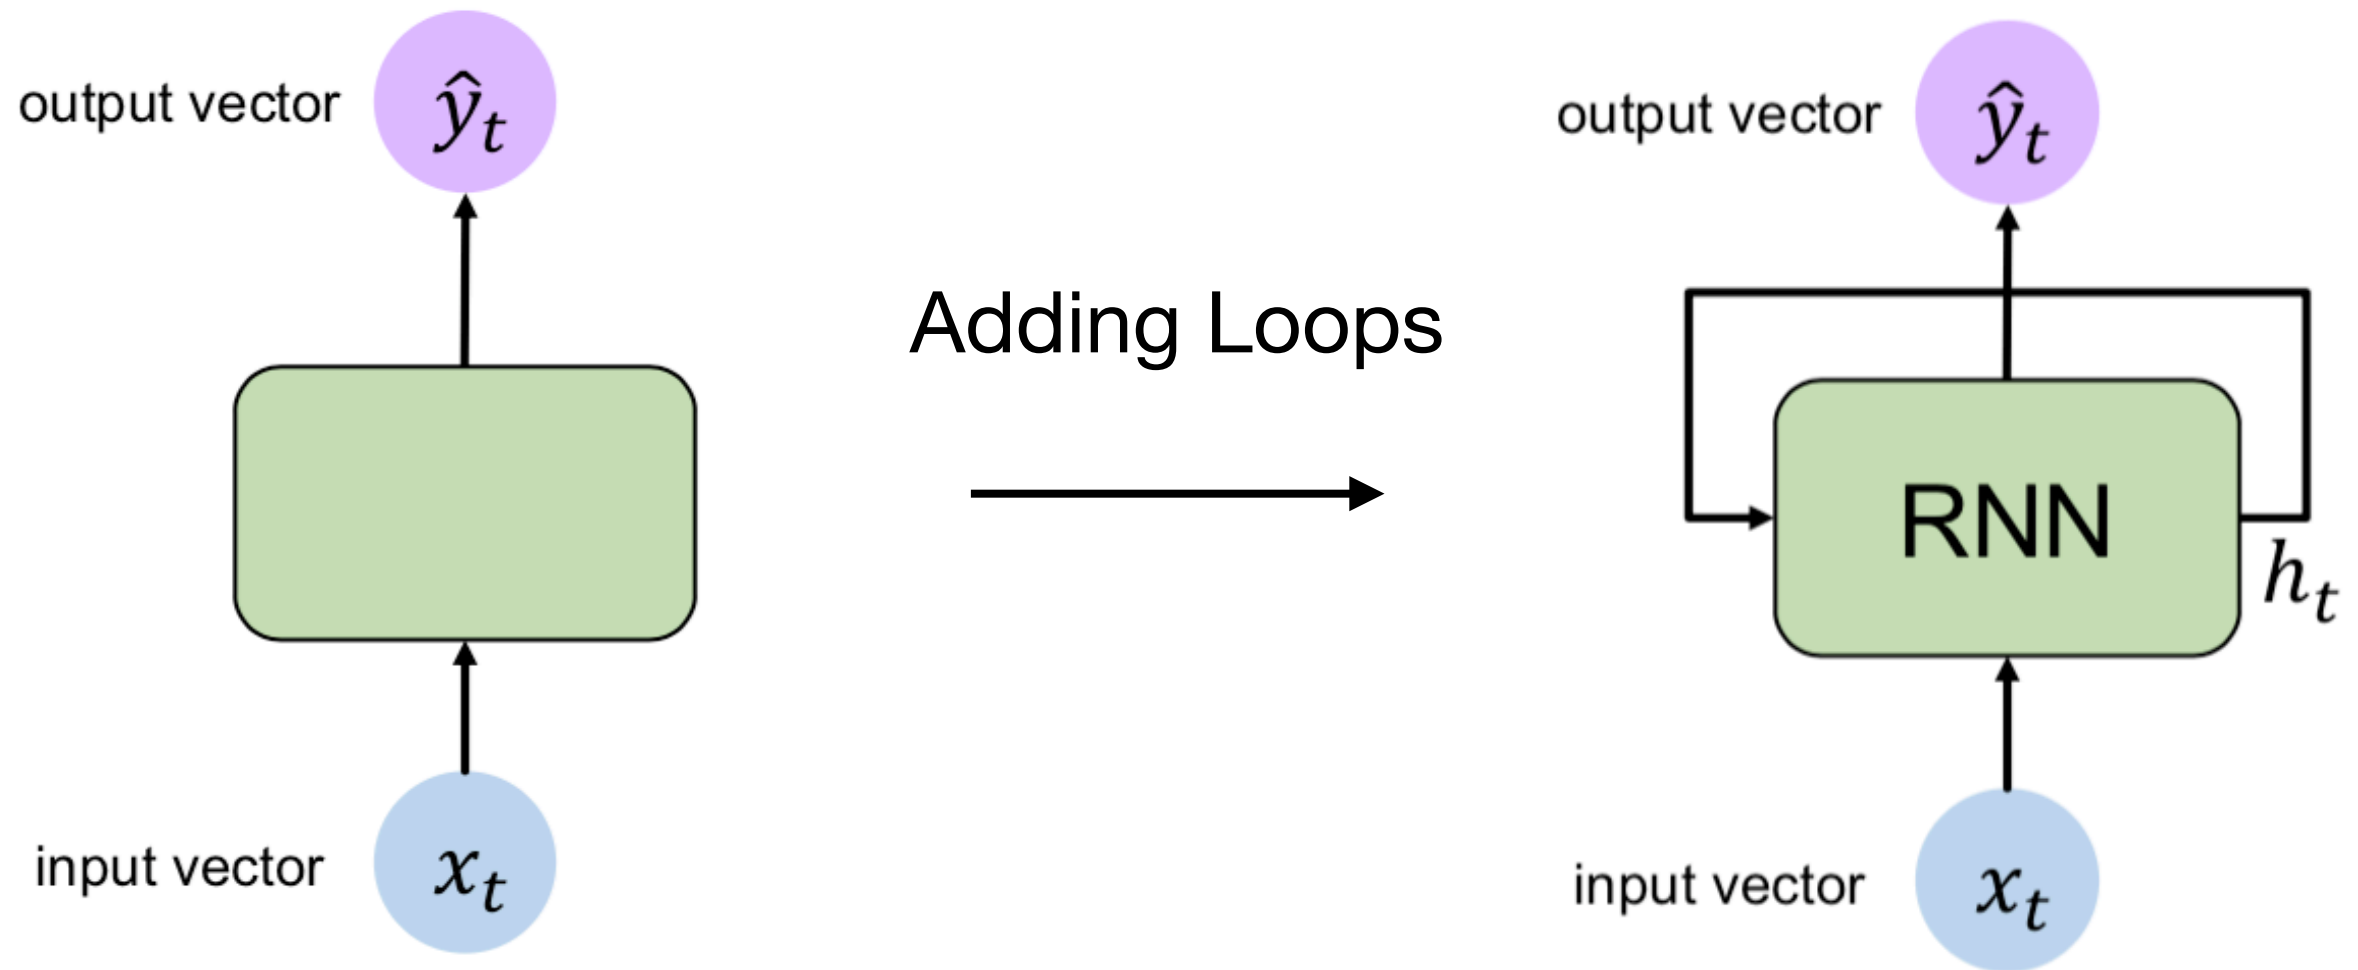
## **Many to May**

For example: Q and A system.

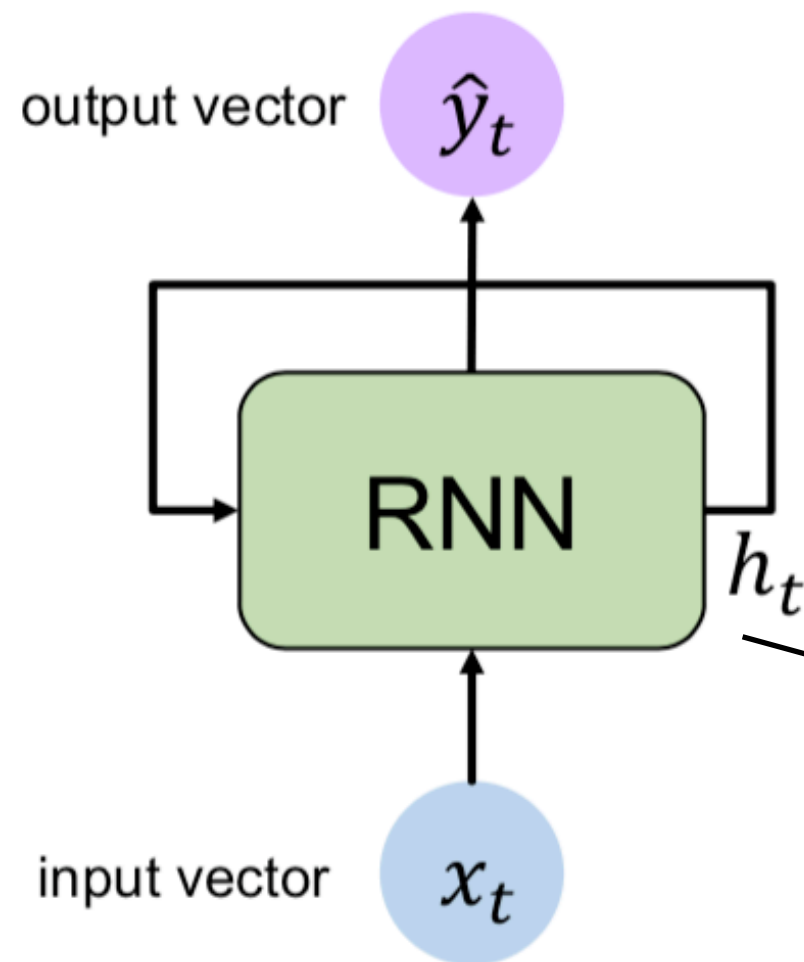
Input: which city is the capital of China?

Output: the capital of China is Beijing.

# Recurrent neural networks



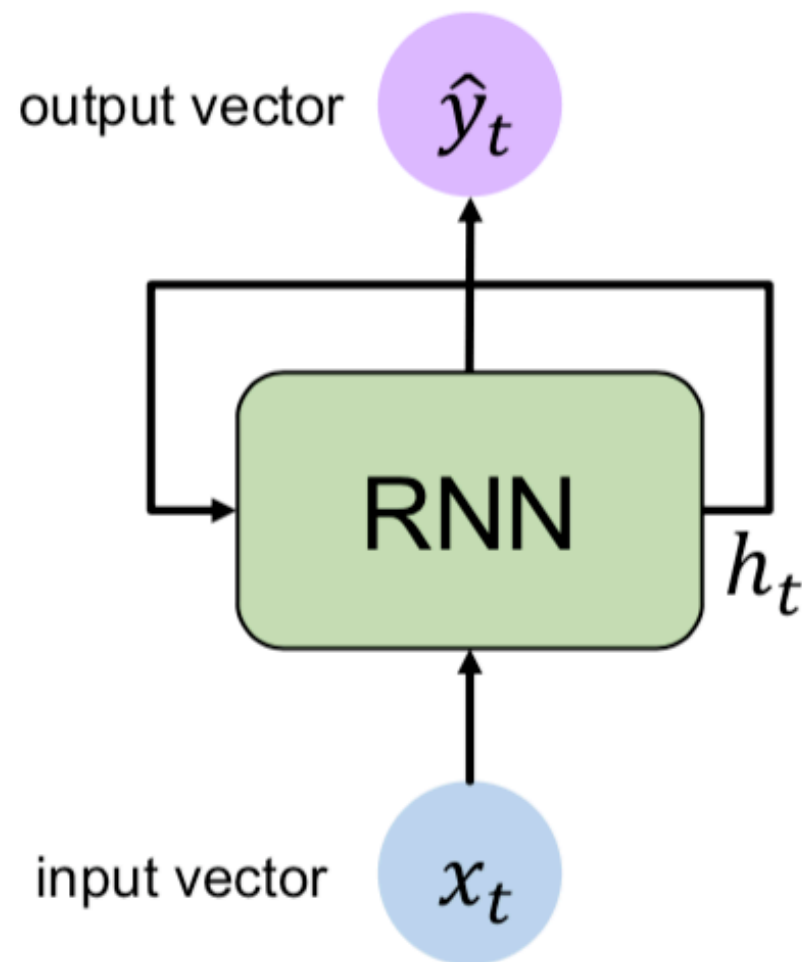
# Recurrent neural networks



$$h_t = f_w(h_{t-1}, x_t)$$

- Recurrent cell
- $h_t$  : hidden state at time t step

# Recurrent neural networks

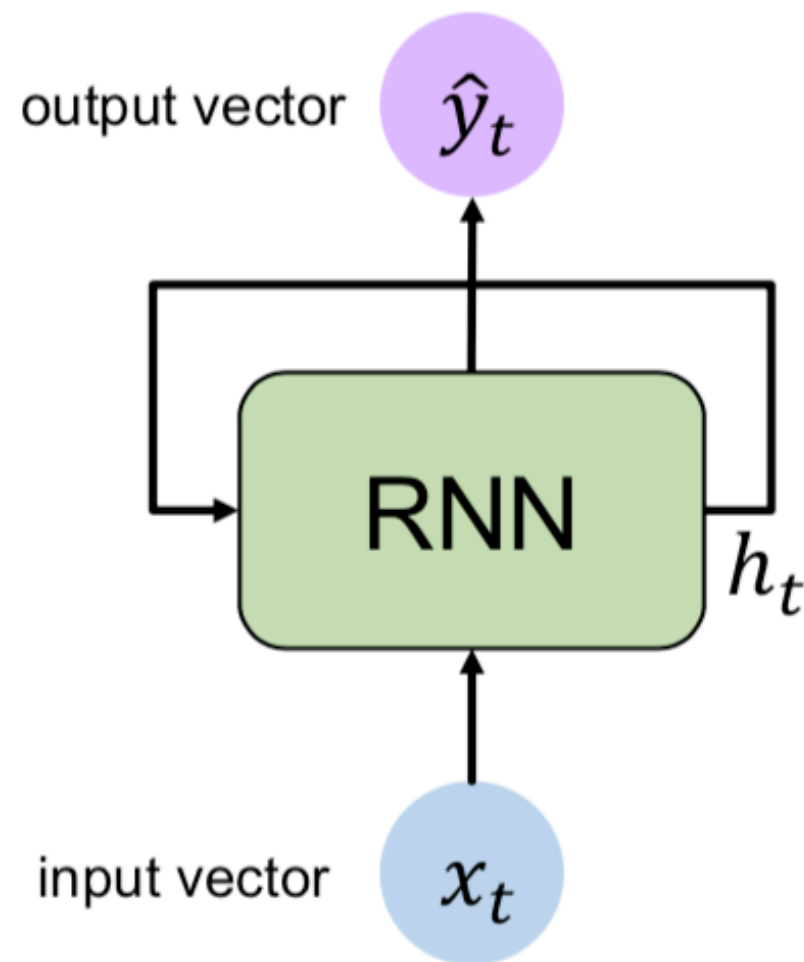


$$\boxed{h_t} = \boxed{f_W}(\boxed{h_{t-1}}, \boxed{x_t})$$

new state      function parameterized by  $W$       old state      input vector at time step  $t$

$h_t$  is the “memory” of the network. It’s calculated based on the previous hidden state and the input at the current step.

# Recurrent neural networks



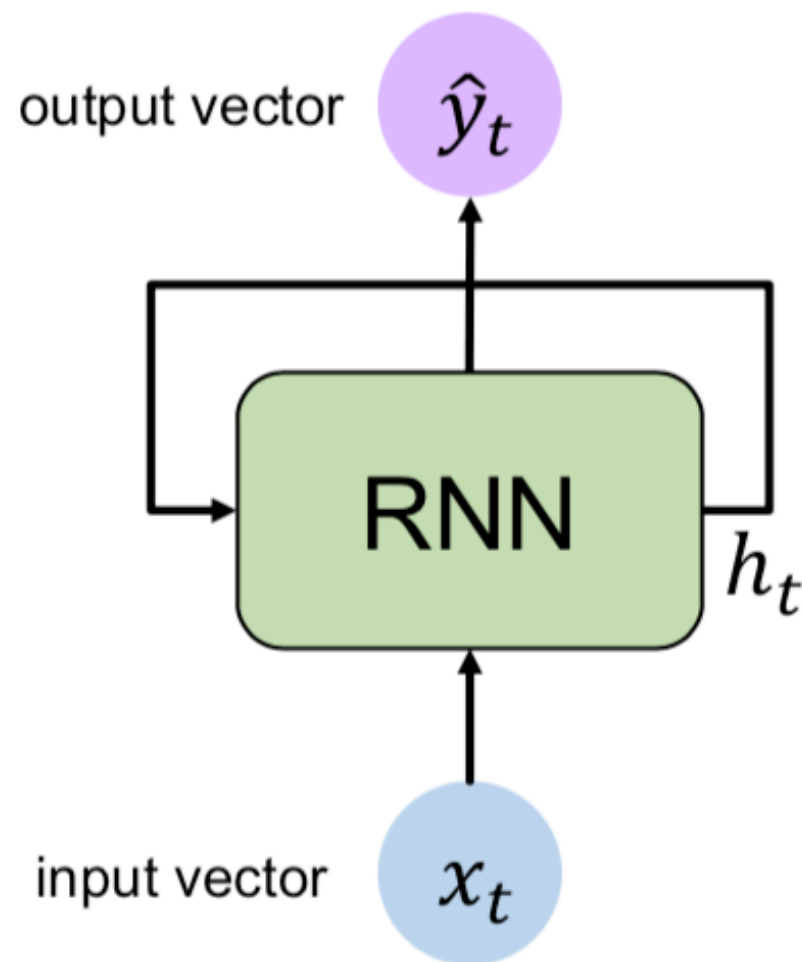
$$\boxed{h_t} = \boxed{f_W}(\boxed{h_{t-1}}, \boxed{x_t})$$

new state      function parameterized by  $W$       old state      input vector at time step  $t$

The same function  $f_W$  and set of parameters  $W$  are used at every time step



# Recurrent neural networks

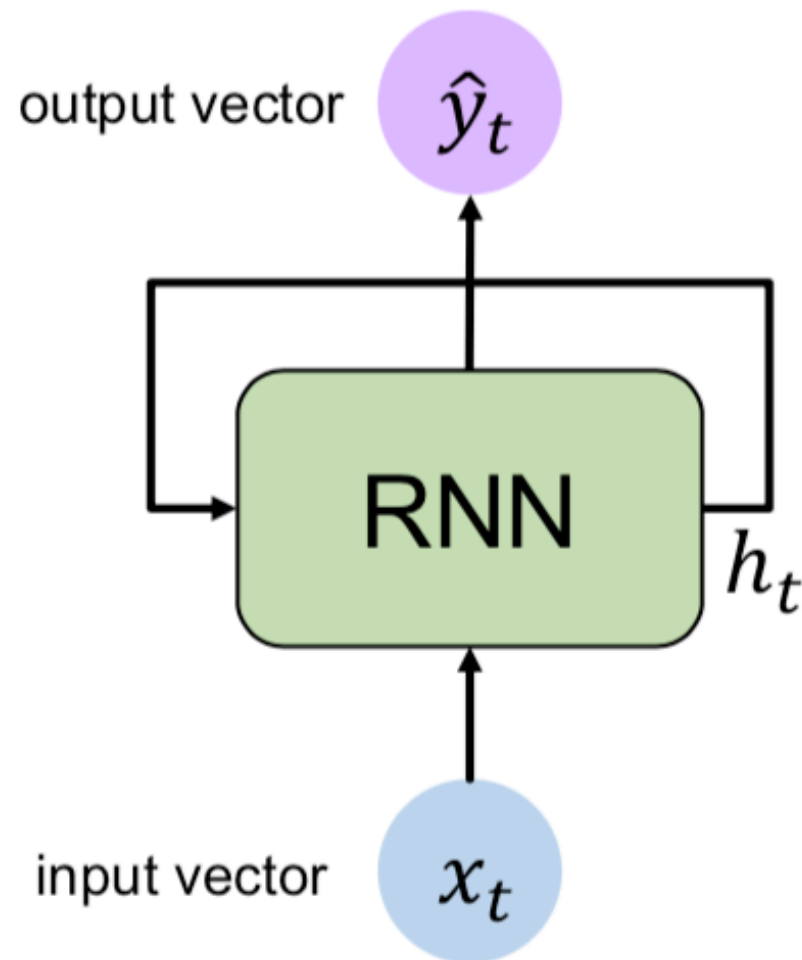


$$h_t = \tanh(\mathbf{W}_{hh}h_{t-1} + \mathbf{W}_{xh}x_t)$$

Two weight matrices

- Input vector  $x_t$
- Previous state  $h_{t-1}$

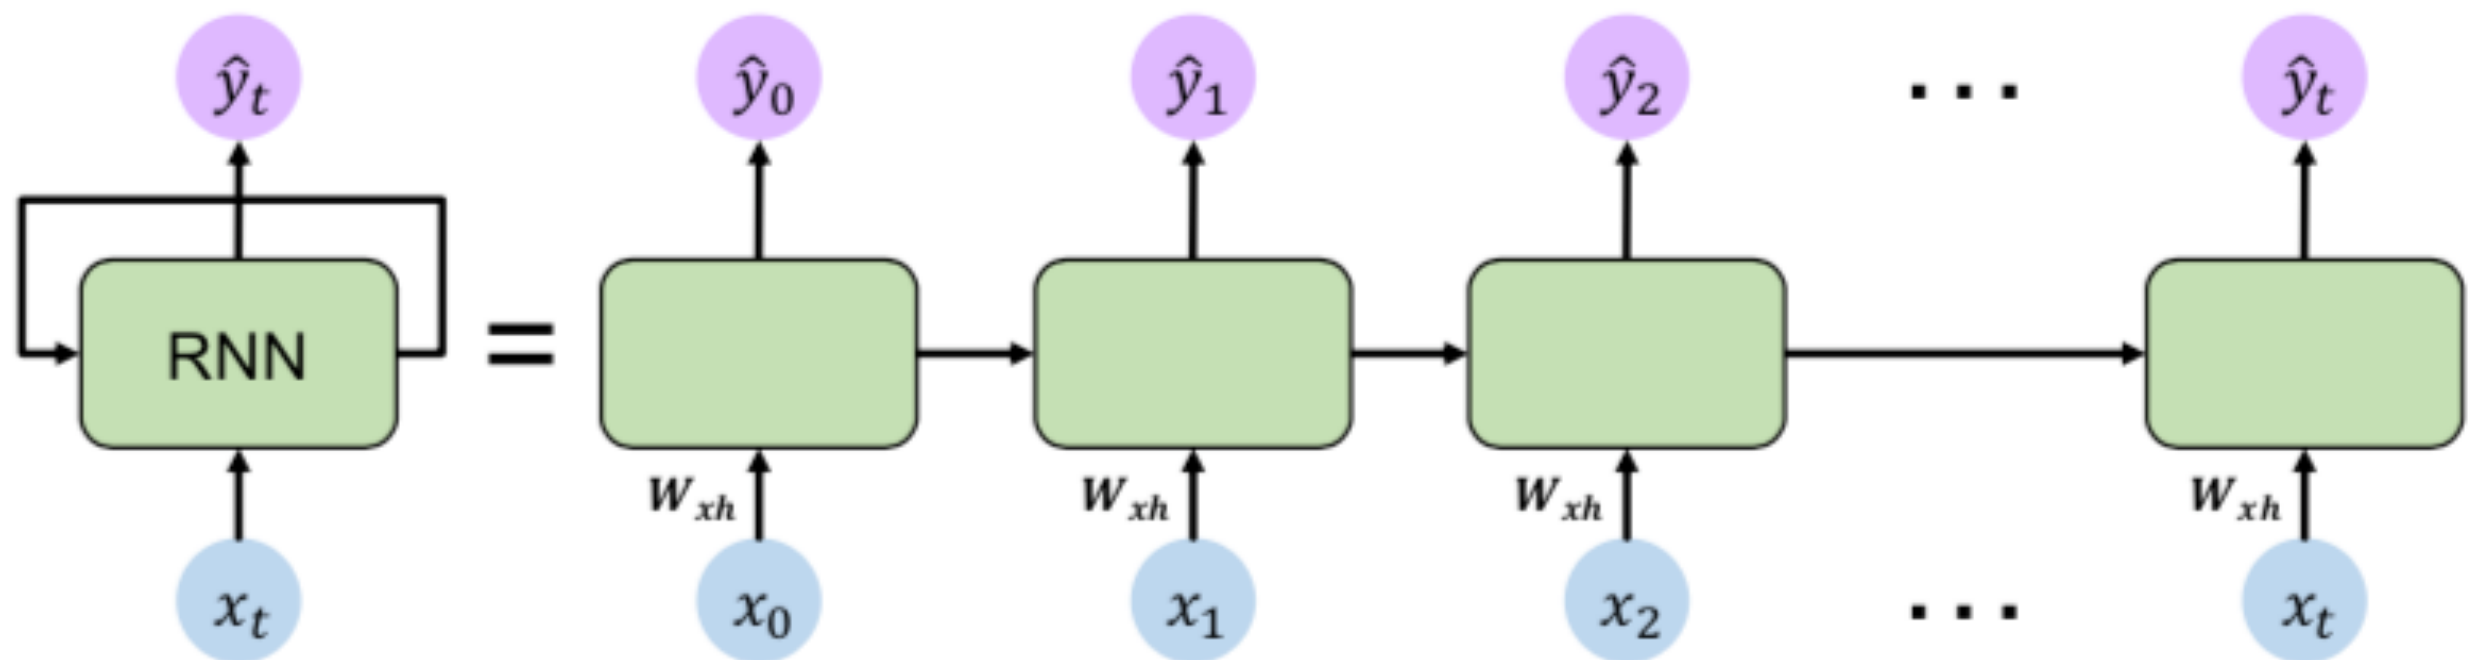
# Recurrent neural networks



$$\hat{y}_t = \mathbf{W}_{hy}h_t$$

$$h_t = \tanh(\mathbf{W}_{hh}h_{t-1} + \mathbf{W}_{xh}x_t)$$

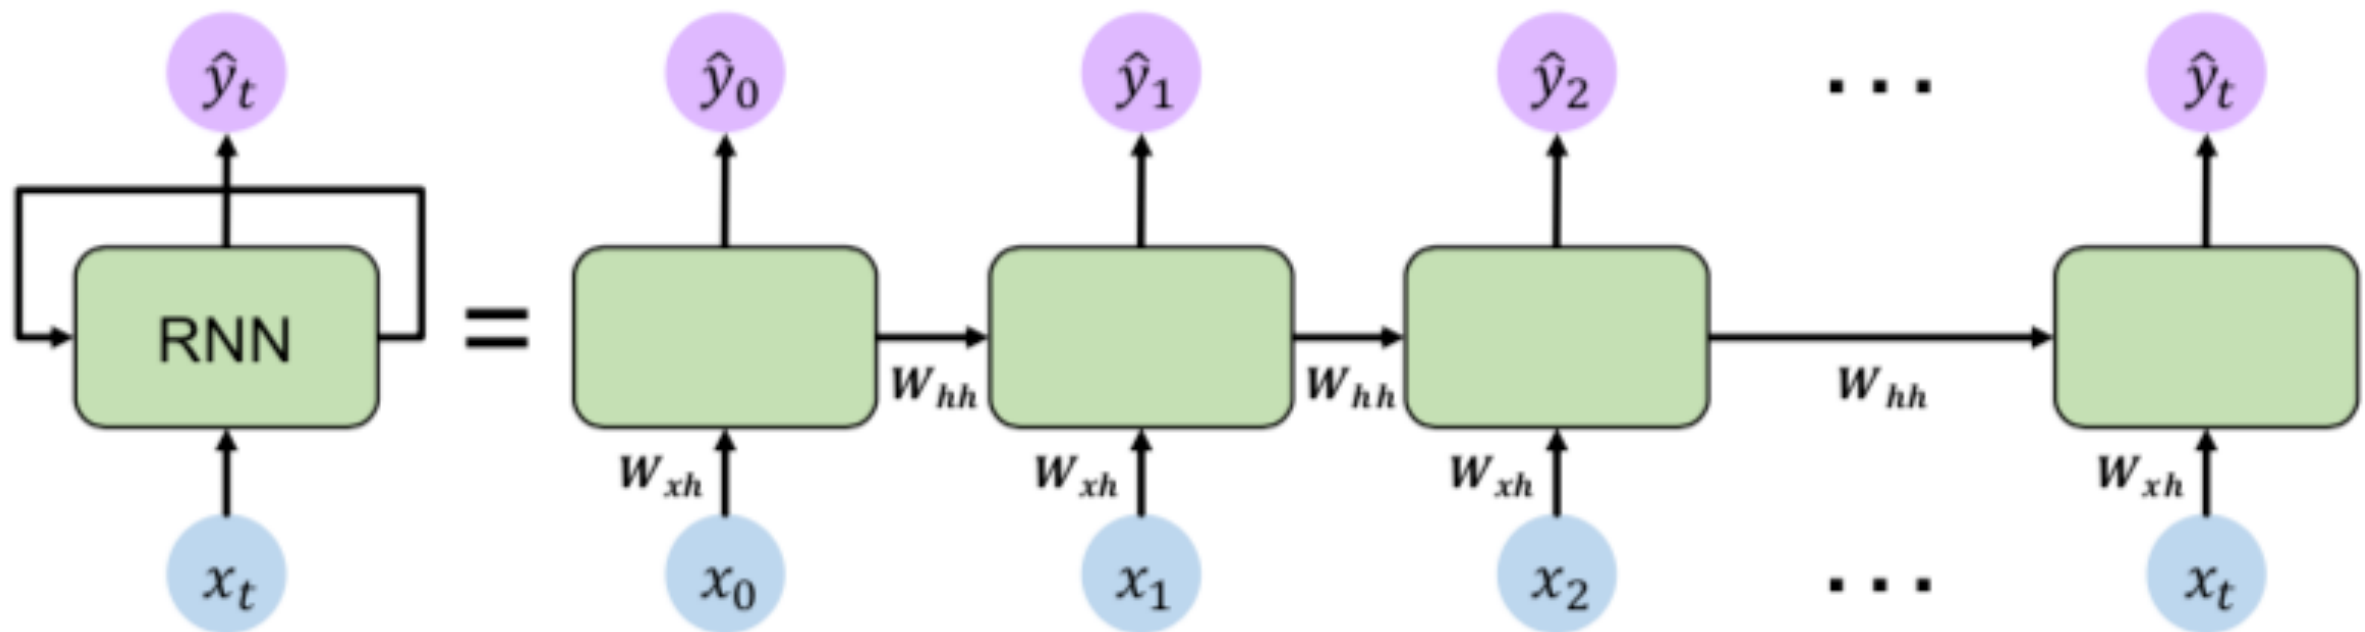
# RNNs: computational graph across time



$x_0, x_1, x_2 \dots x_t$  are the inputs

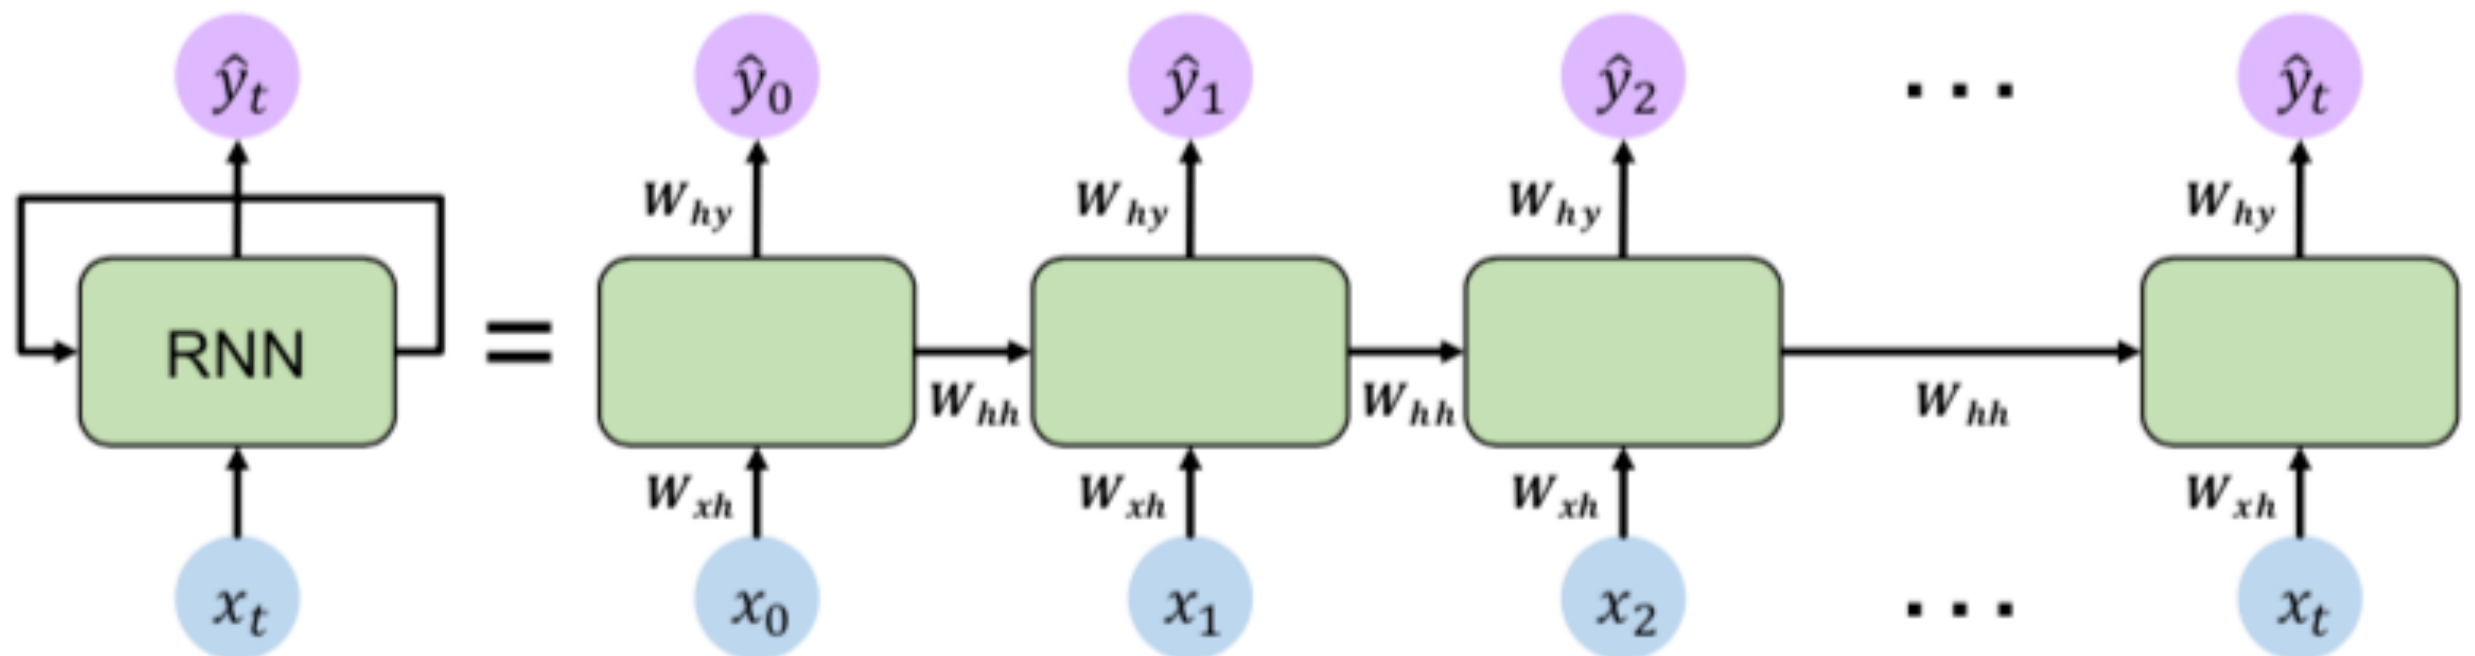
$W_{xh}$  is the weight of inputs

# RNNs: computational graph across time



$W_{hh}$  is the weight of hidden state cell.

# RNNs: computational graph across time



$W_{hy}$  is the weight of outputs.

# Recurrent neural networks

- RNNs are called recurrent because they perform the same task for every element of a sequence, with the output being depended on the previous computations.
- With hidden state cell, they have a “memory” which captures information about what has been calculated so far.
- Unlike a traditional deep neural network, which uses different parameters at each layer, a RNN shares the same parameters ( $W_{hh}$ ,  $W_{xh}$ ,  $W_{hy}$ ) across all steps. This means we are performing the same task at each step, just with different inputs. This greatly reduces the total number of parameters we need to learn.



# The Problems Of RNNs

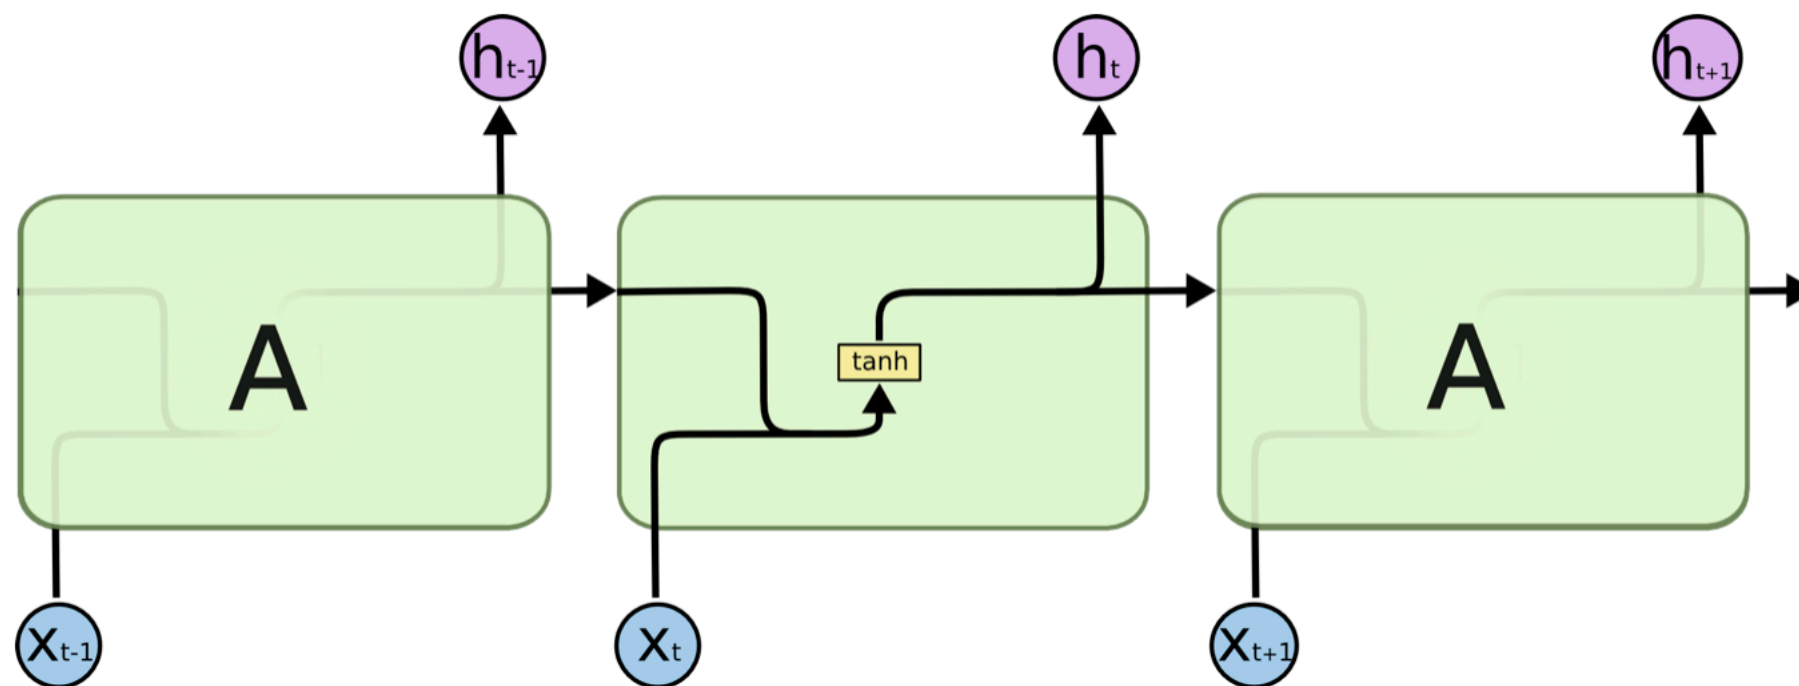
- RNNs can learn to use the past information when the gap between the relevant information and the place that it's needed is small.
- But when the gap between the relevant information is large, RNNs are unable to learn to connect the relevant information!

# LSTM networks

Long Short Term Memory networks (LSTMs) are a special kind of RNN, are a special kind of RNN, capable of learning **long-term dependencies**.

# Standard RNNs VS LSTMs

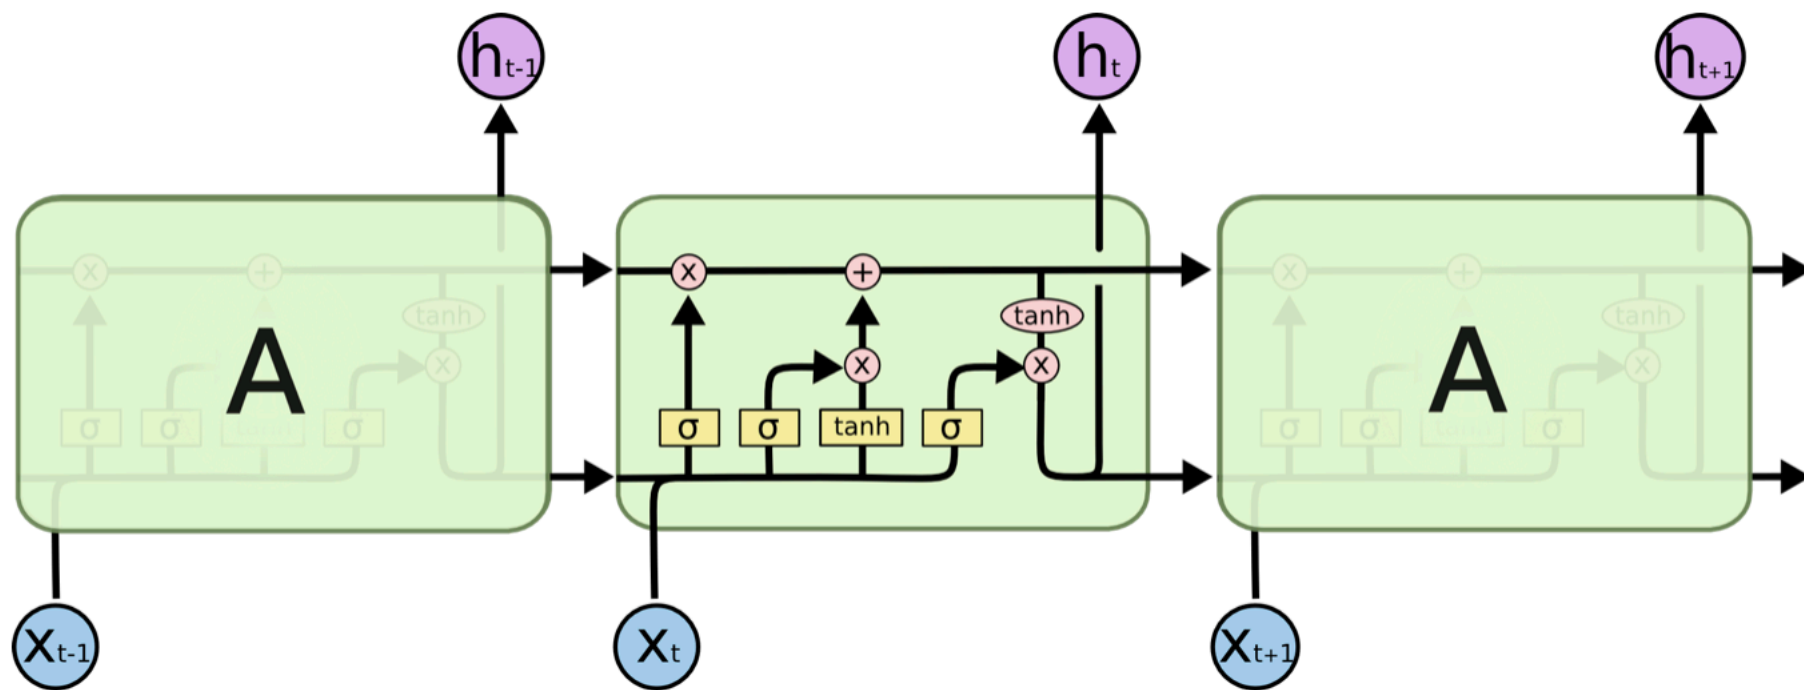
In standard RNNs, this repeating module will have a very simple structure, such as a single tanh layer.



The repeating module in a standard RNN contains a single layer.

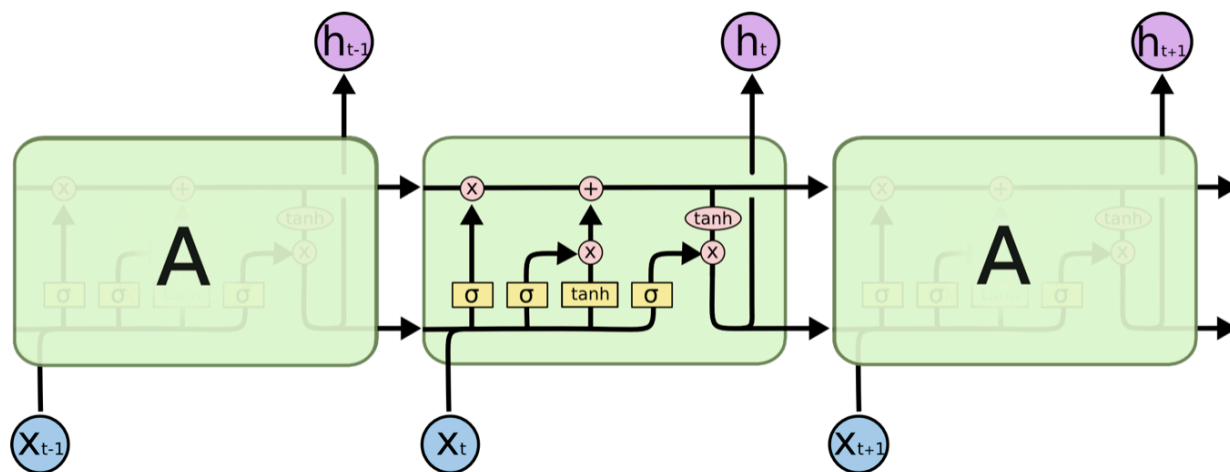
# Standard RNNs VS LSTMs

- In LSTMs, the repeating module has four neural network layers, instead of one, interacting in a very special way.

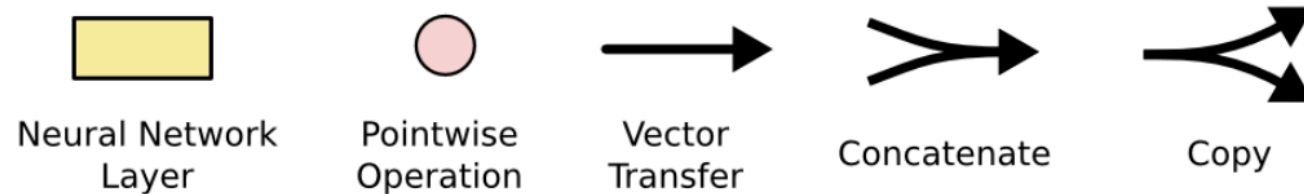


The repeating module in an LSTM contains four interacting layers.

# Symbols Of LSTMs



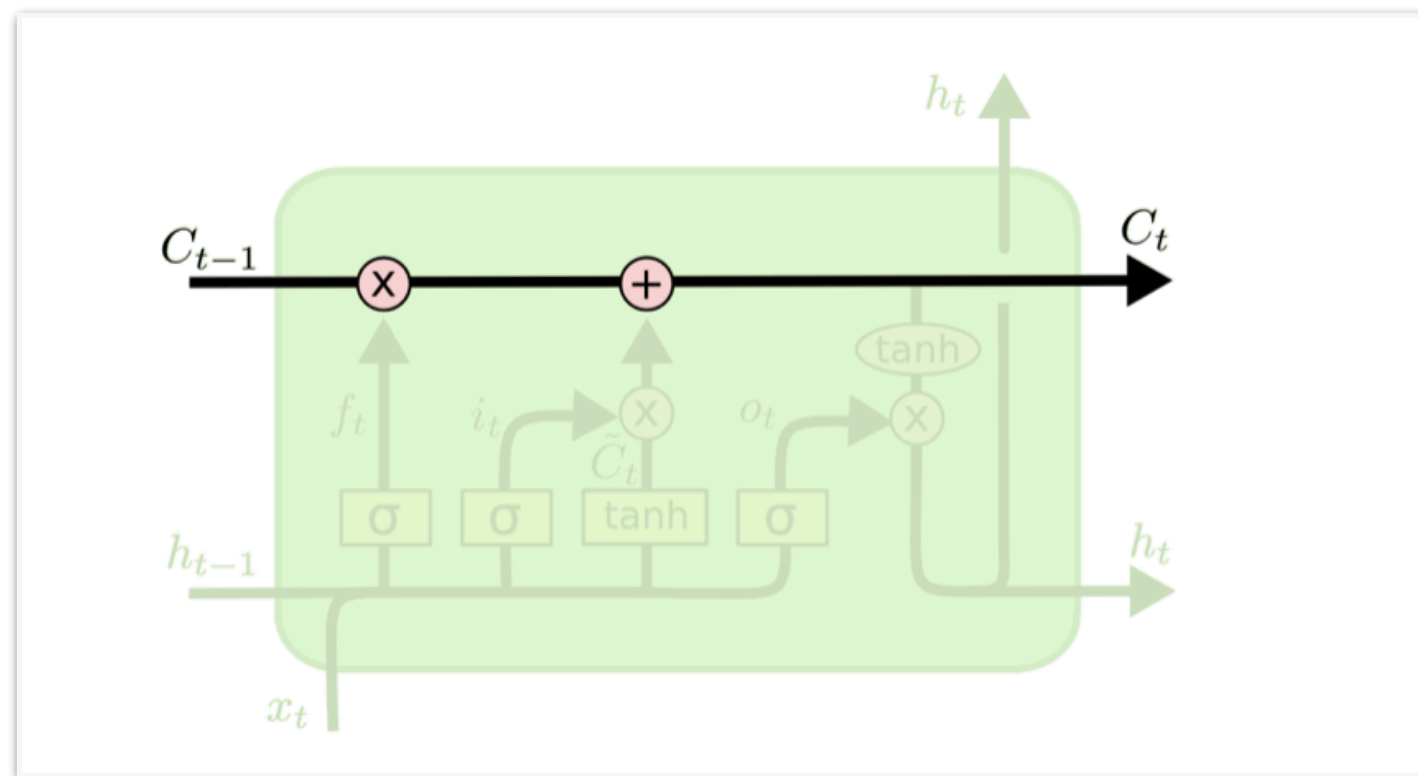
The repeating module in an LSTM contains four interacting layers.



- Each line carries an entire vector, from the output of one node to the inputs of others.
- The pink circles represent pointwise operations, like vector addition, while the yellow boxes are learned neural network layers.
- Lines merging denote concatenation, while a line forking denote its content being copied and the copies going to different locations.

# The Core Idea Behind LSTMs

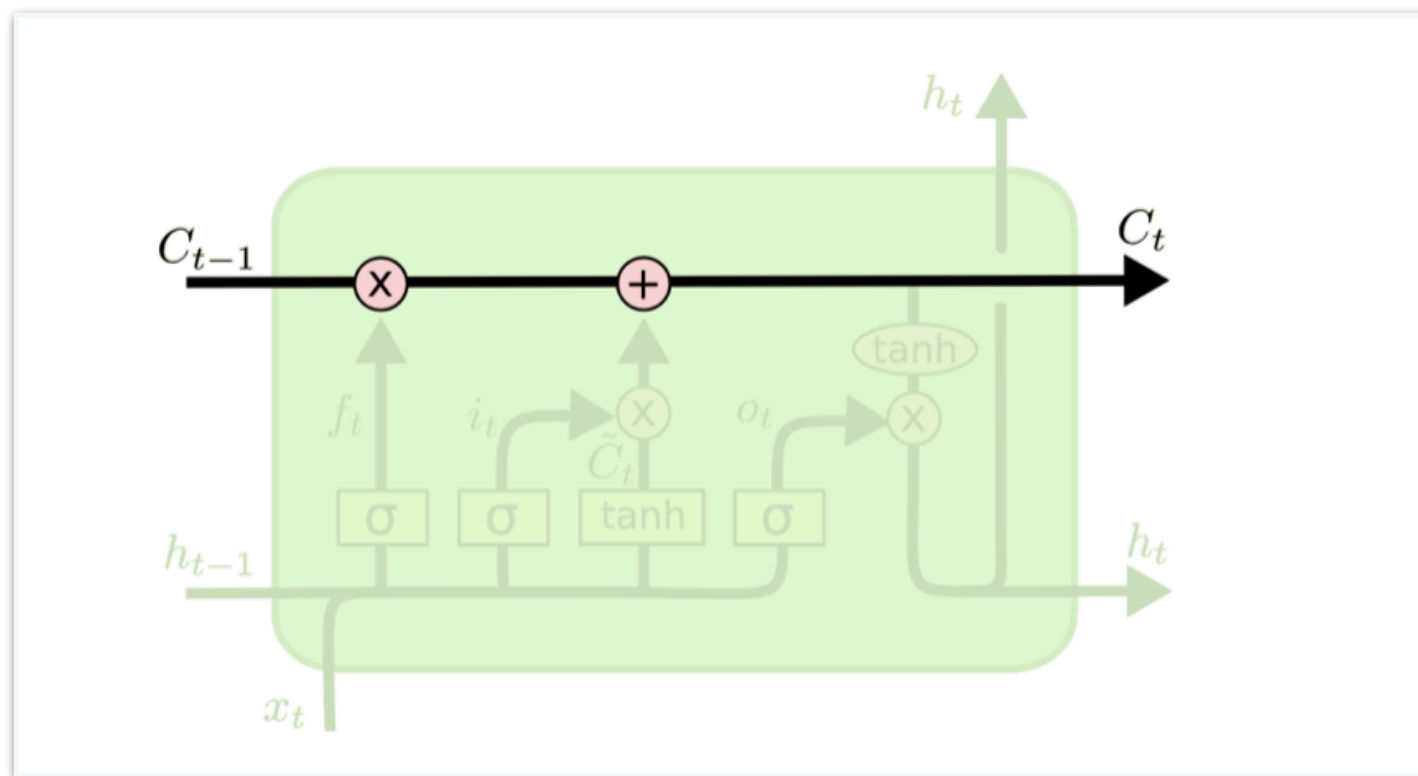
- **Cell state** is the horizontal line running through the top of the diagram, which is the **key** to LSTMs





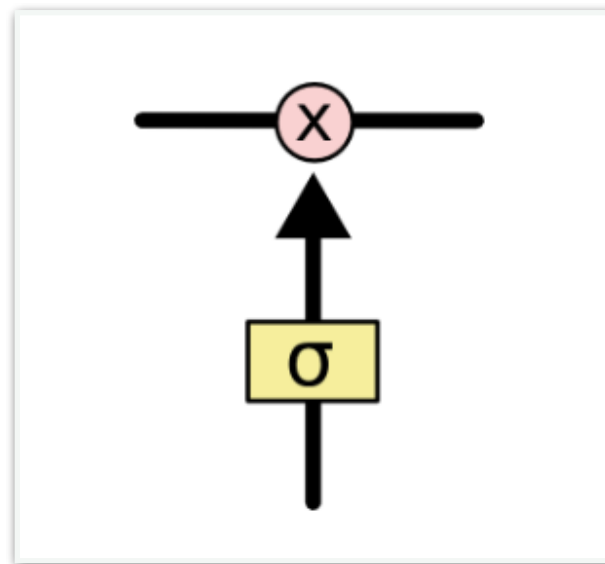
# The Core Idea Behind LSTMs

- Cell state runs straight down the entire chain, with only some minor linear interactions. It's very easy for information to just flow along it unchanged.



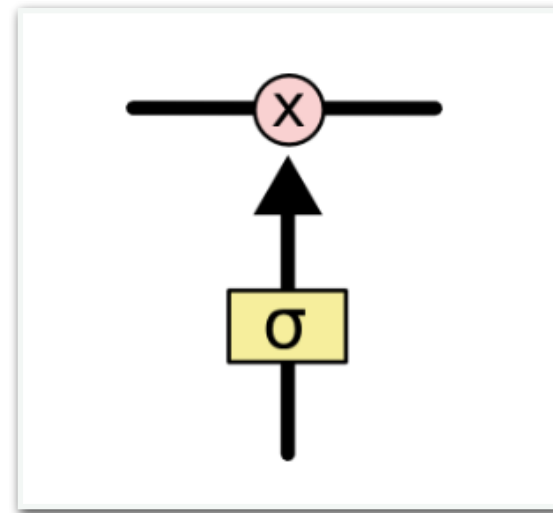
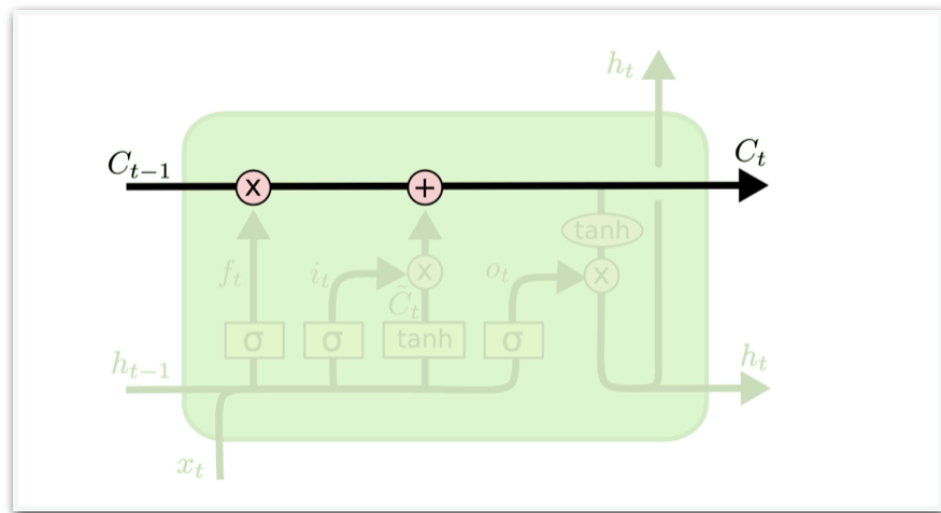
# The Core Idea Behind LSTMs

- **Gates** are a way to optionally let information through. They are composed out of a **sigmoid** neural net layer and a pointwise **multiplication** operation.



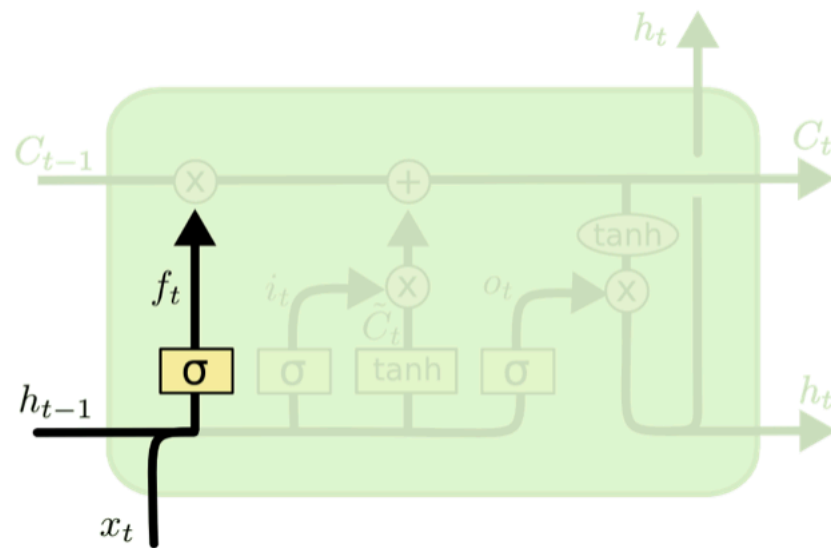
# The Core Idea Behind LSTMs

- The LSTM use gates remove or add information to the cell state.



# Step-by-Step LSTM

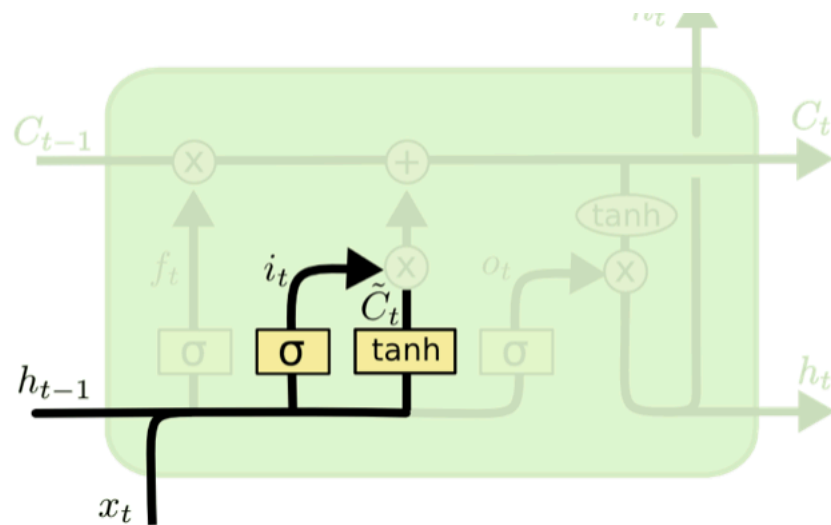
- The first step in our LSTM is to decide what information we're going to throw away from the cell state.



$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$

# Step-by-Step LSTM

- The second step is to decide what new information we're going to store in the cell state.

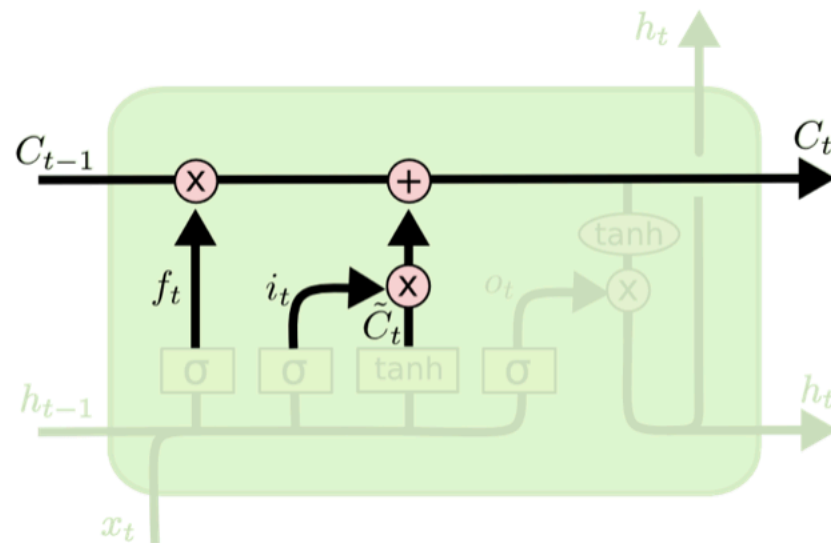


$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

# Step-by-Step LSTM

- The third step is update the old cell state into the new cell state

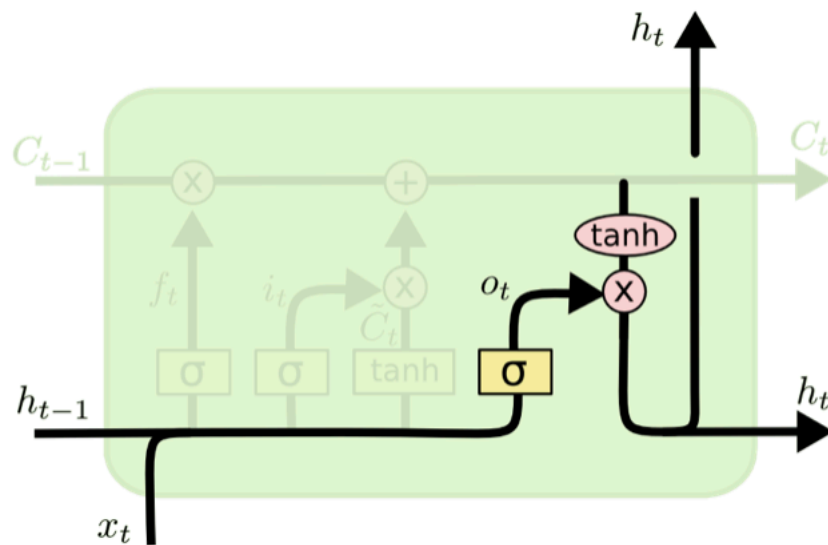


$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$



# Step-by-Step LSTM

- Finally, we need to decide what we're going to output



$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$

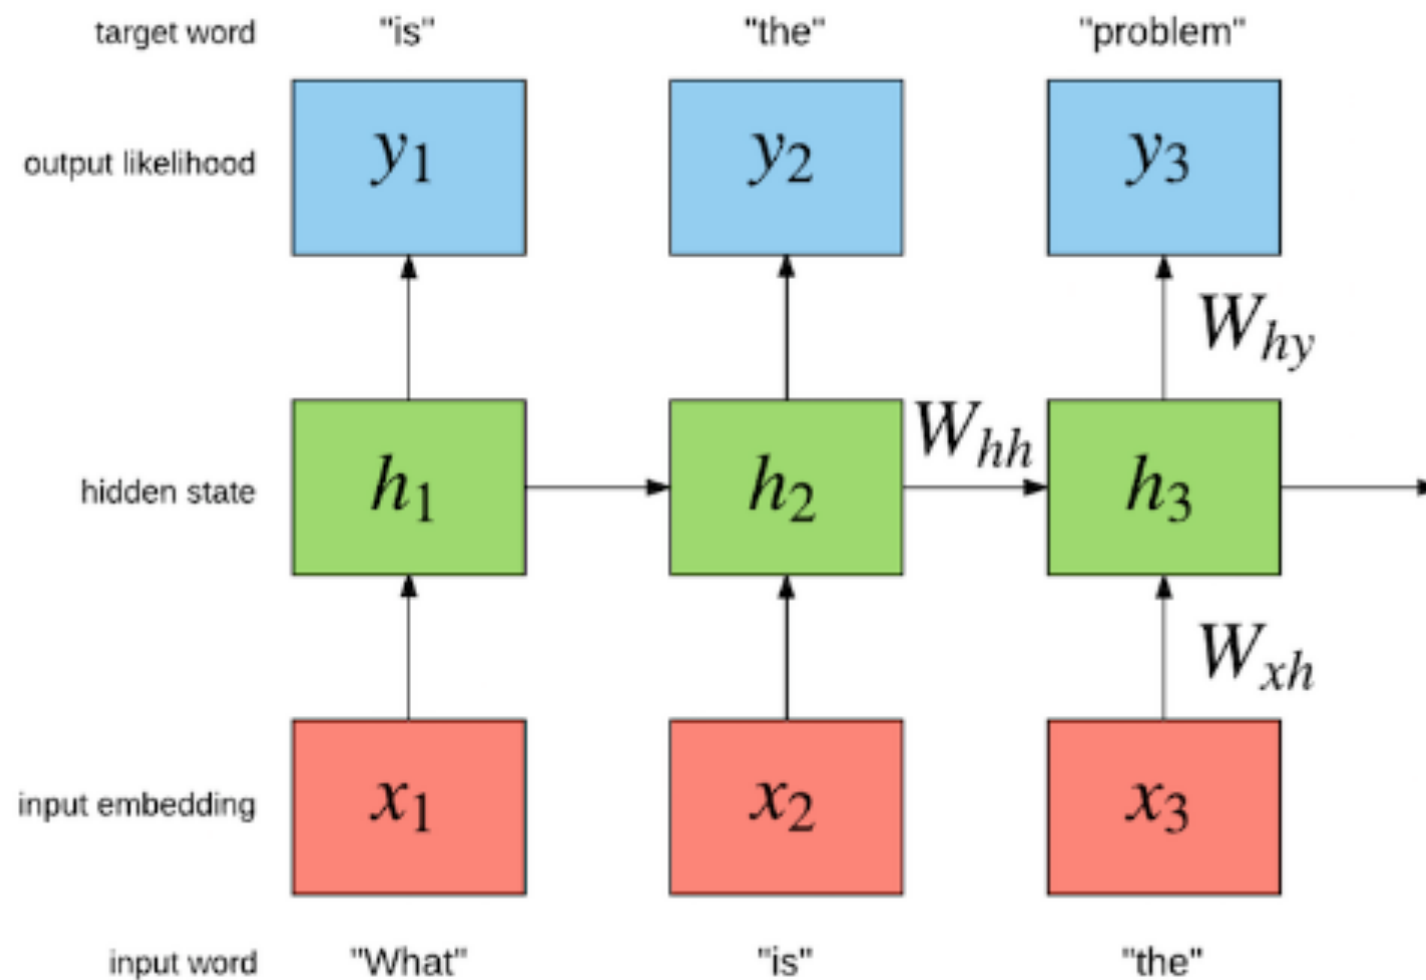
$$h_t = o_t * \tanh(C_t)$$

# LSTMs: key concepts

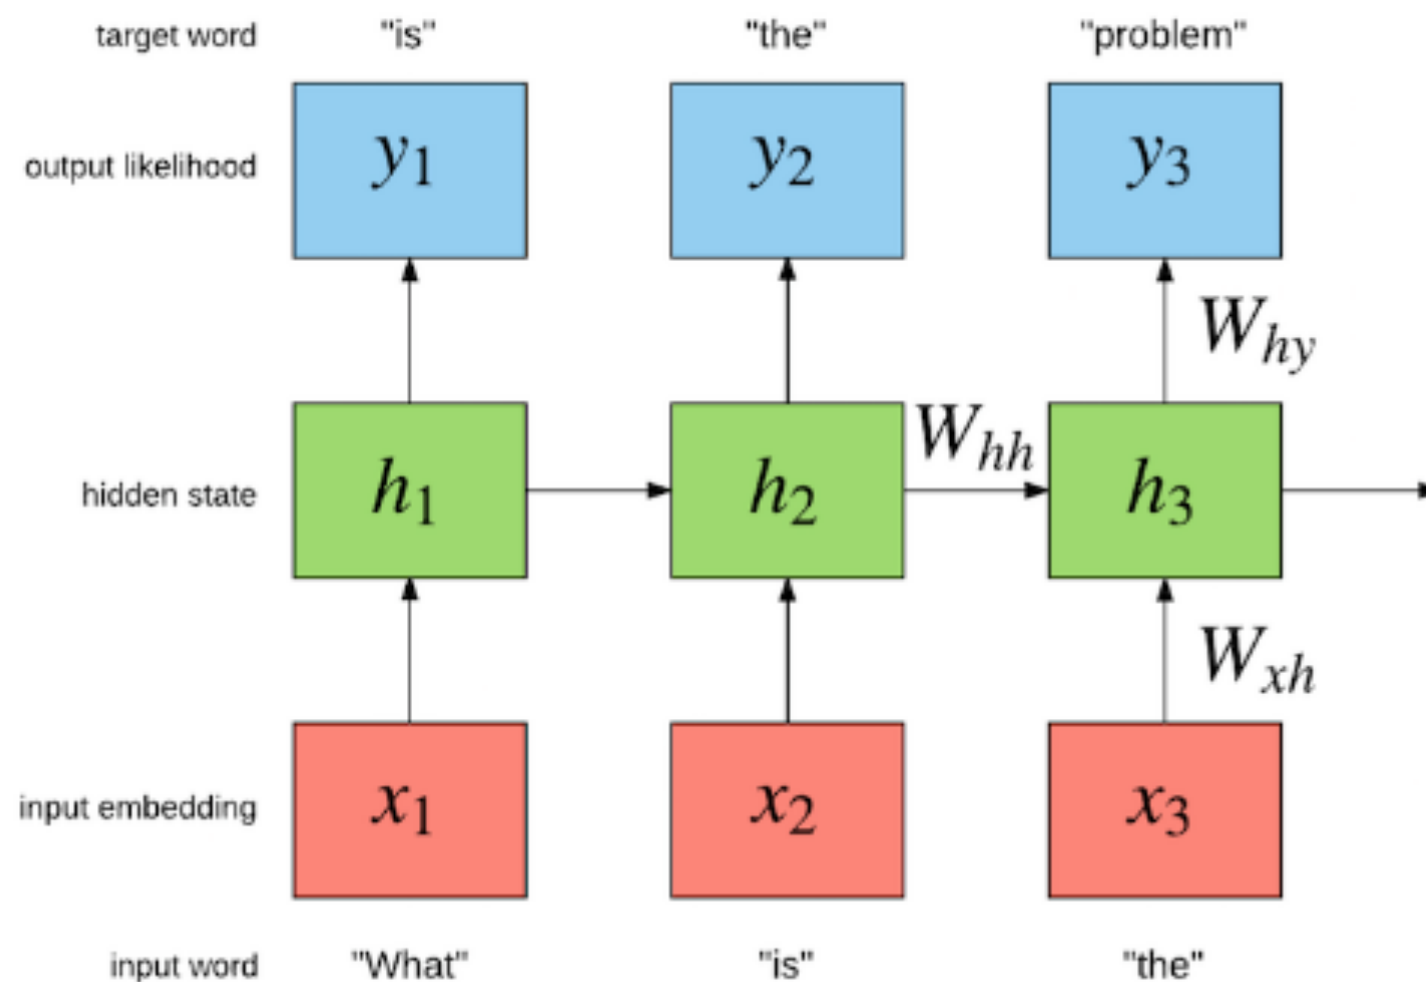
1. Maintain a separate cell state from what is outputted
2. Use **gates** to control the flow of information
  - Forget gate gets rid of irrelevant information
  - Selectively update cell state
  - Output gate returns a filtered version of the cell state

# Application of RNNs

- Language Model
- Language Model



# Language Model

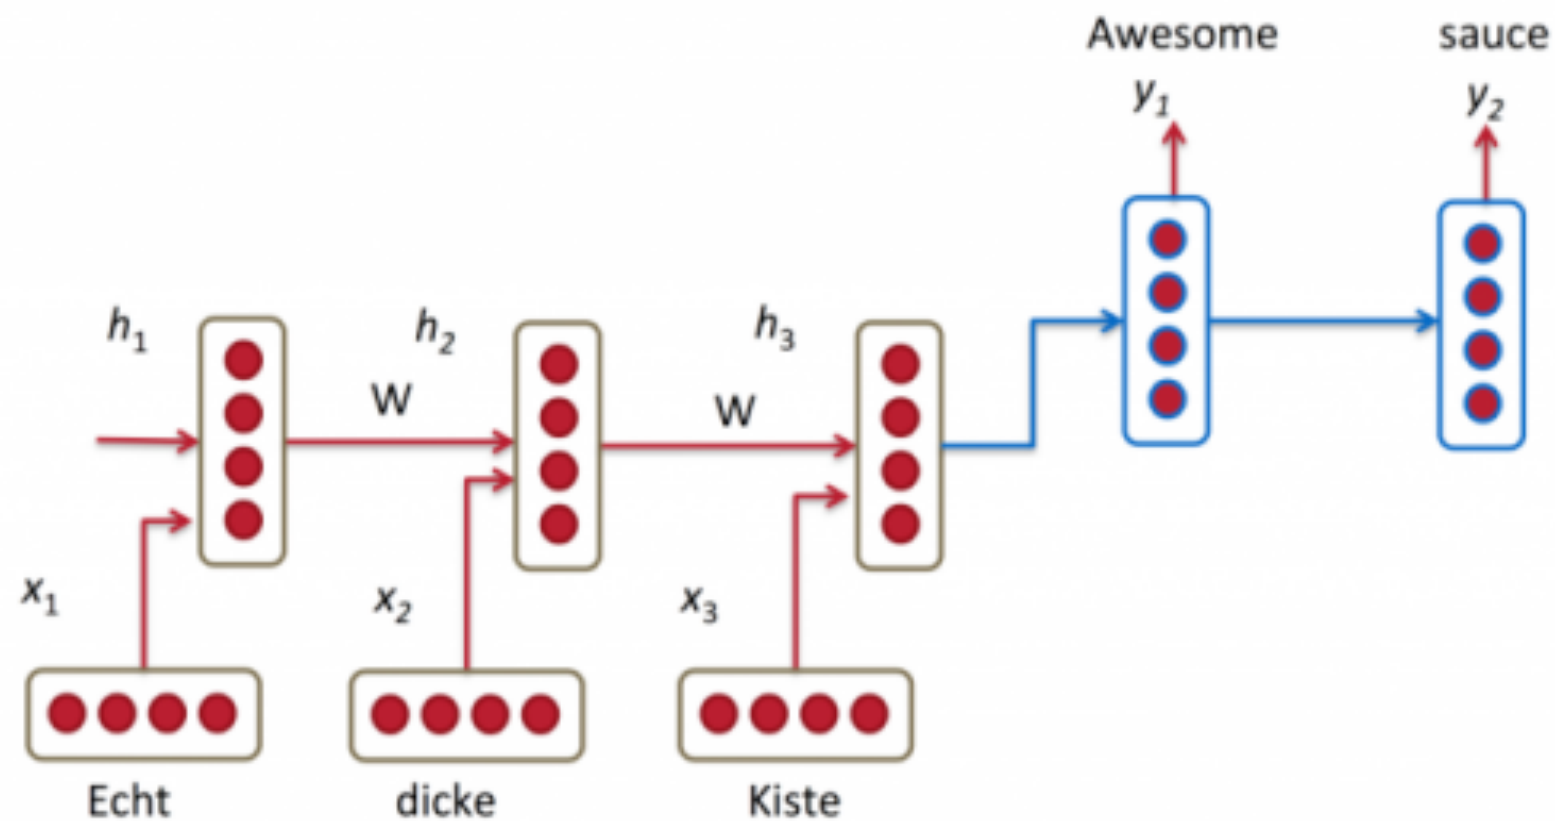


# Writing a poem

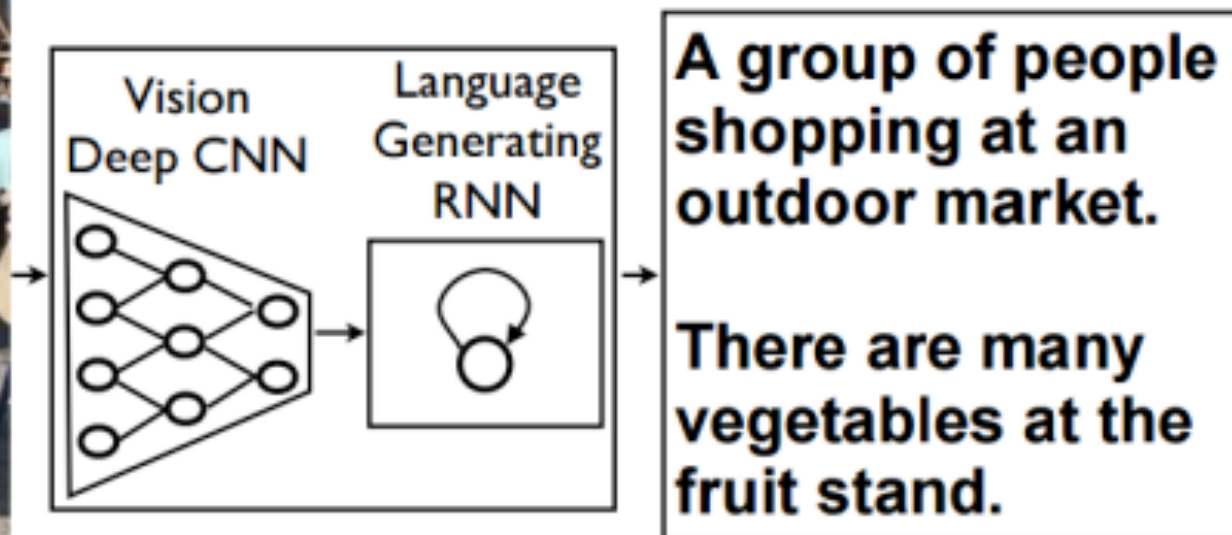
<p>白鹭窥鱼立， Egrets stood, peeping fishes. 青山照水开。 Water was still, reflecting mountains. 夜来风不动， The wind went down by nightfall, 明月见楼台。 as the moon came up by the tower.</p>	<p>满怀风月一枝春， Budding branches are full of romance. 未见梅花亦可人。 Plum blossoms are invisible but adorable. 不为东风无此客， With the east wind comes Spring. 世间何处是前身。 Where on earth do I come from?</p>
--	--

# Machine Translate

- An RNN for coding
- Another RNN for decoding

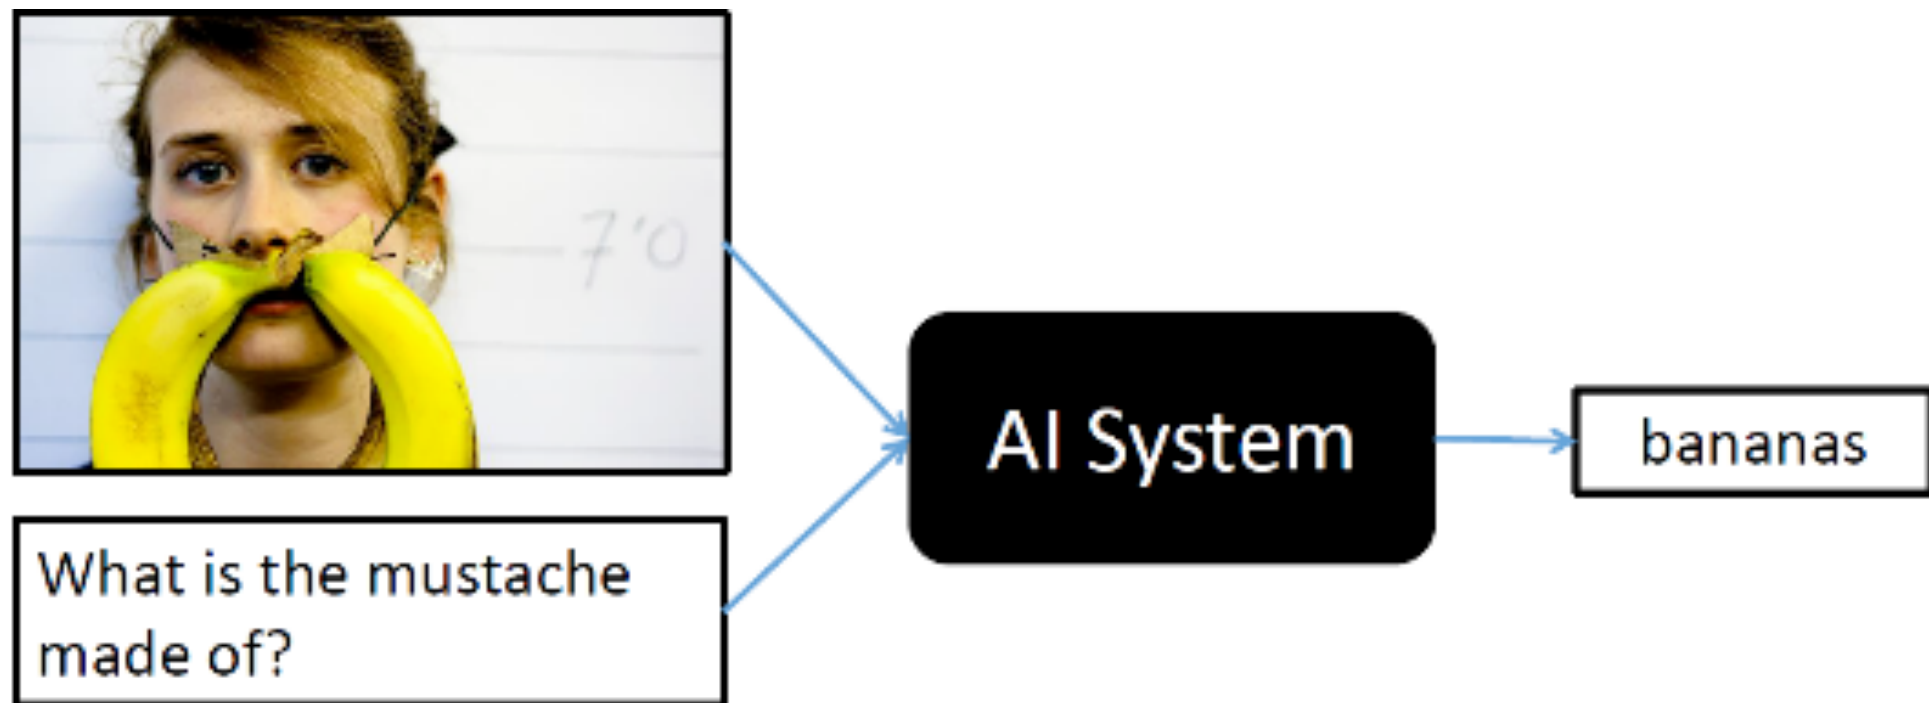


# Talking about pictures



# Visual Question Answering (VQA)

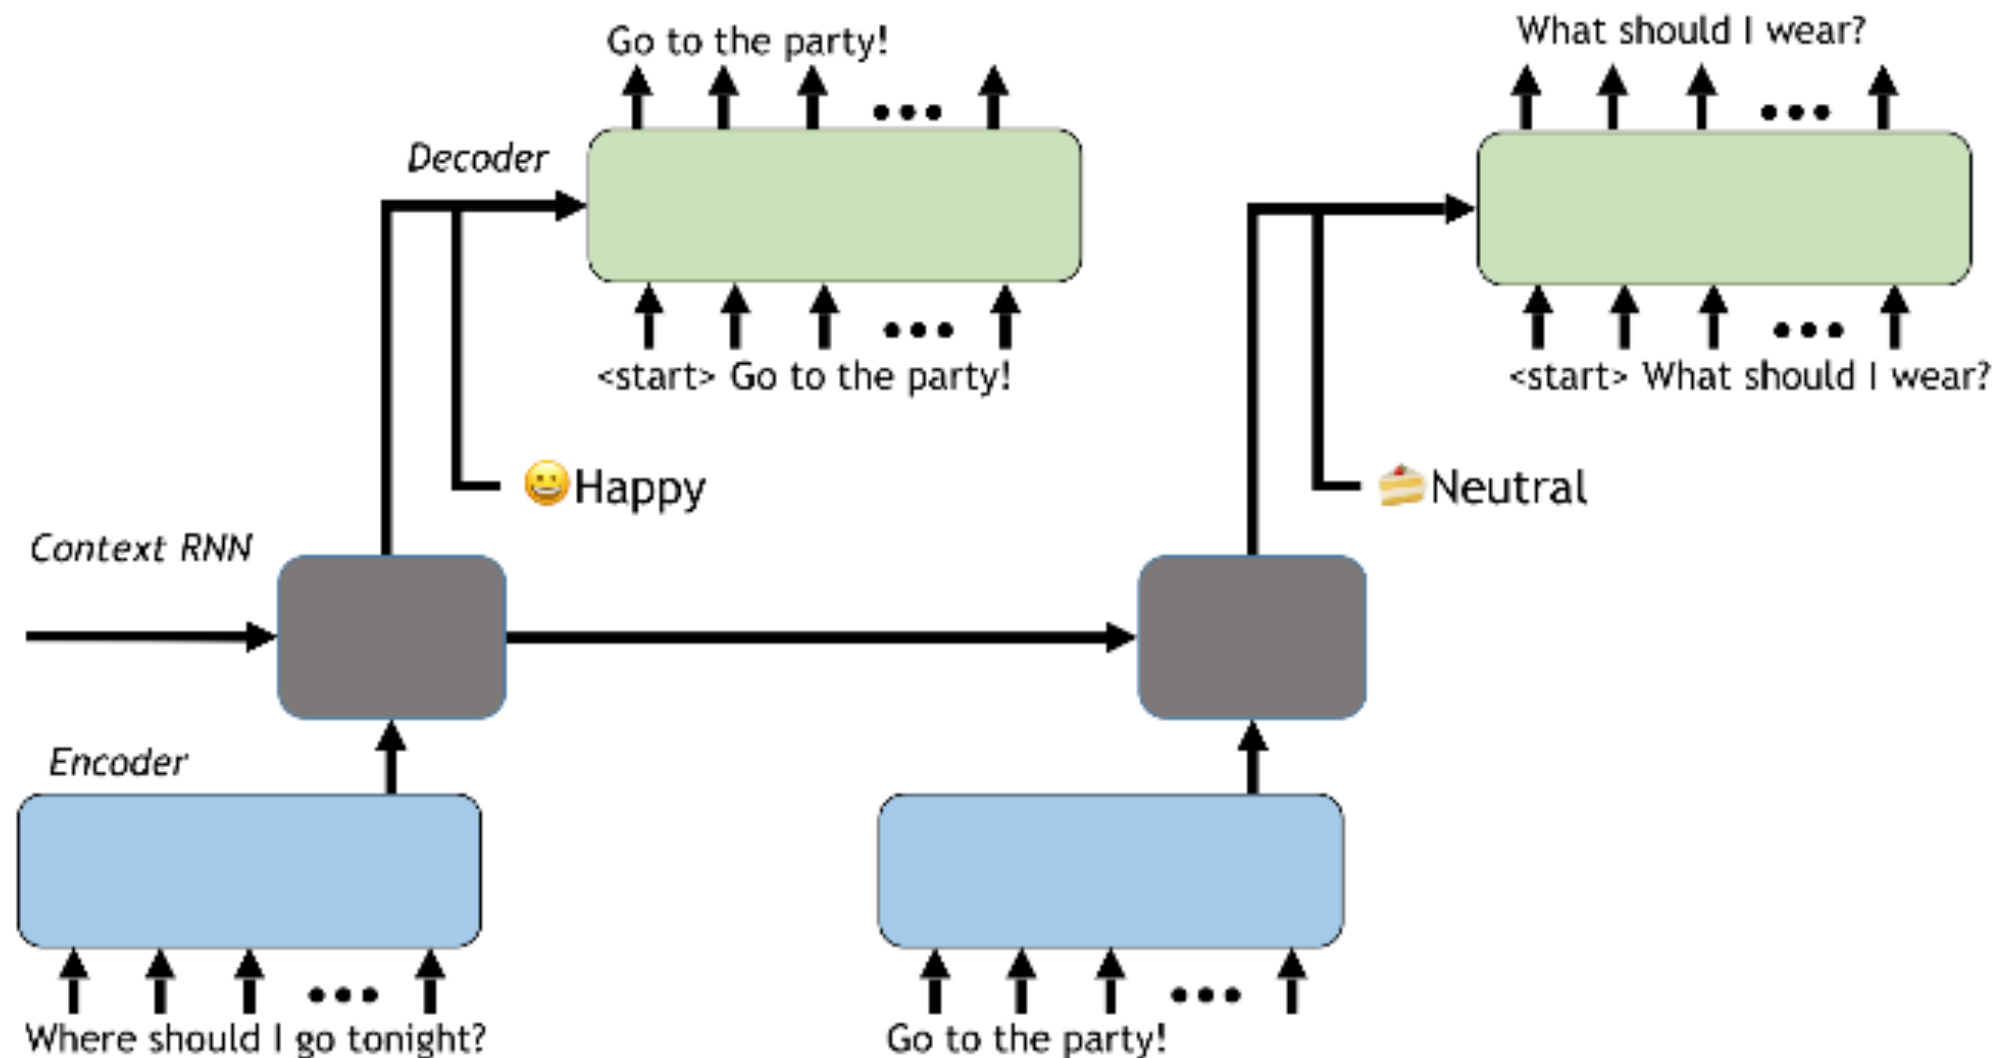
VQA: Given an image and a natural language question about the image, the task is to provide an accurate natural language answer





# Dialogue system

<https://github.com/lukalabs/cakechat>



**Thank you!**